EXP NO: 5   **INSTALLATION OF HIVE**

$wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz


$tar xvzf apache-hive-3.1.2-bin.tar.gz
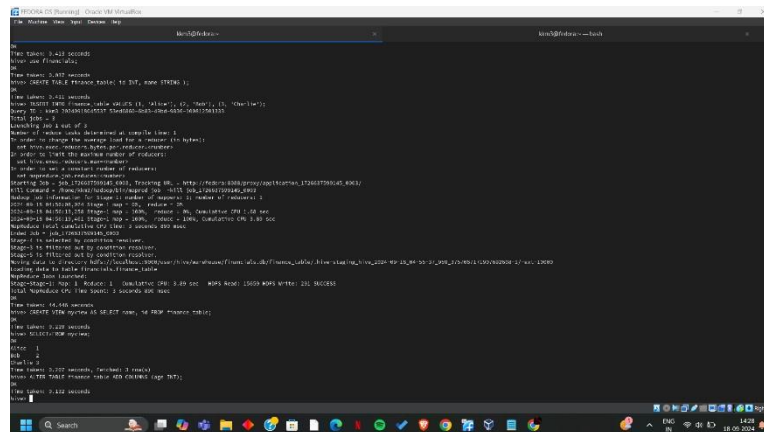
$mv apache-hive-3.1.2-bin hive

$nano ~/.bashrc

$source ~/.bashrc

$cd hive

$cp conf/hive-default.xml.template conf/hive-site.xml

$nano conf/hive-site.xml
$start-all.sh

$ jps

$hdfs dfs -mkdir /user

$hdfs dfs -mkdir /user/hive

$hdfs dfs -mkdir /user/hive/warehouse

$hdfs dfs -mkdir /user/tmp

$hdfs dfs -chmod g+w /user/tmp

$hdfs dfs -chmod g+w /user/hive/warehouse

$cd $HIVE_HOME/bin

$schematool -initSchema –dbType derby

$hive

>show databases;

> CREATE DATABASE financials;

> use financials;

> CREATE TABLE finance_table(id INT, name STRING);

> INSERT INTO finance_tableVALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');

```
hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240910011318_759e97f4-dbcc-4020-a345-37b6aac0dbbf
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1725938931910_0001, Tracking URL = http://kali:8088/proxy/application_1725938931910_0001/
Kill Command = /home/hadoop/hadoop/bin/mapred job  -kill job_1725938931910_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-09-10 01:14:33,477 Stage-1 map = 0%,  reduce = 0%
2024-09-10 01:15:17,392 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 20.45 sec
2024-09-10 01:15:43,645 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 33.61 sec
MapReduce Total cumulative CPU time: 33 seconds 610 msec
Ended Job = job_1725938931910_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-10_01-13-18_205_31425232155846852-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 33.61 sec   HDFS Read: 15668 HDFS Write: 291 SUCCESS
Total MapReduce CPU Time Spent: 33 seconds 610 msec
OK
Time taken: 151.732 seconds
```

> CREATE VIEW myview AS SELECT name, id FROM finance_table;

```
hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;
OK
Time taken: 0.82 seconds
```

> SELECT * FROM myview;

```
hive> SELECT * FROM myview;
OK
Alice   1
Bob     2
Charlie 3
Time taken: 0.716 seconds, Fetched: 3 row(s)
hive>
```

> DESCRIBE finance_table;

```
hive> DESCRIBE finance_table;
OK
id                      int
name                    string
Time taken: 0.348 seconds, Fetched: 2 row(s)
```

> ALTER TABLE finance_table ADD COLUMNS (age INT);

>quit;

```
hive> ALTER TABLE finance_table ADD COLUMNS (age INT);
OK
Time taken: 0.457 seconds
hive> quit;
```